


Chapter 4: Threads

adapted from Silberschatz, Galvin, Gagne

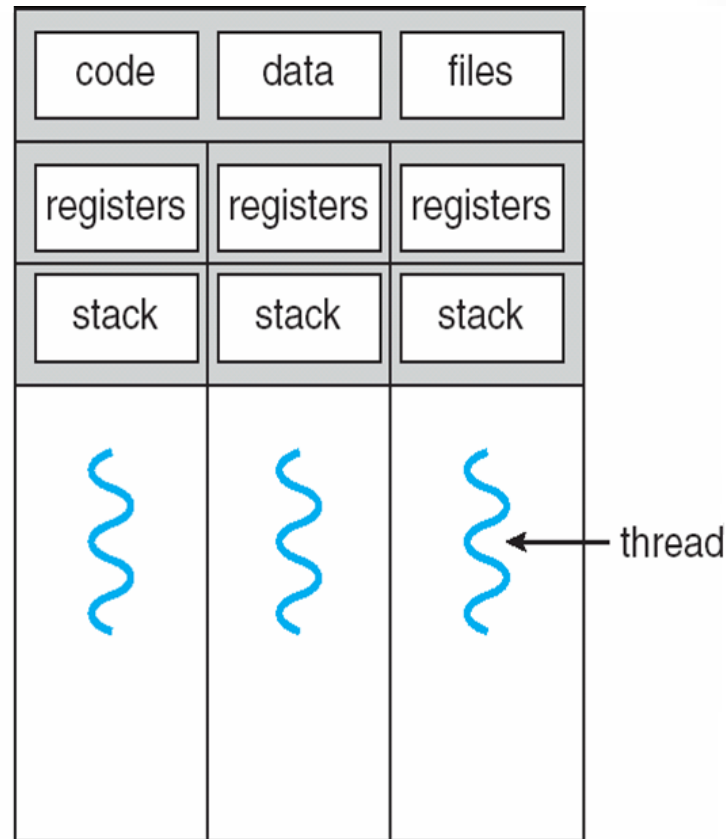
Chapter 4: Threads

- Overview
- Multithreading Models
- Thread Libraries
- Threading Issues

Single and Multithreaded Processes

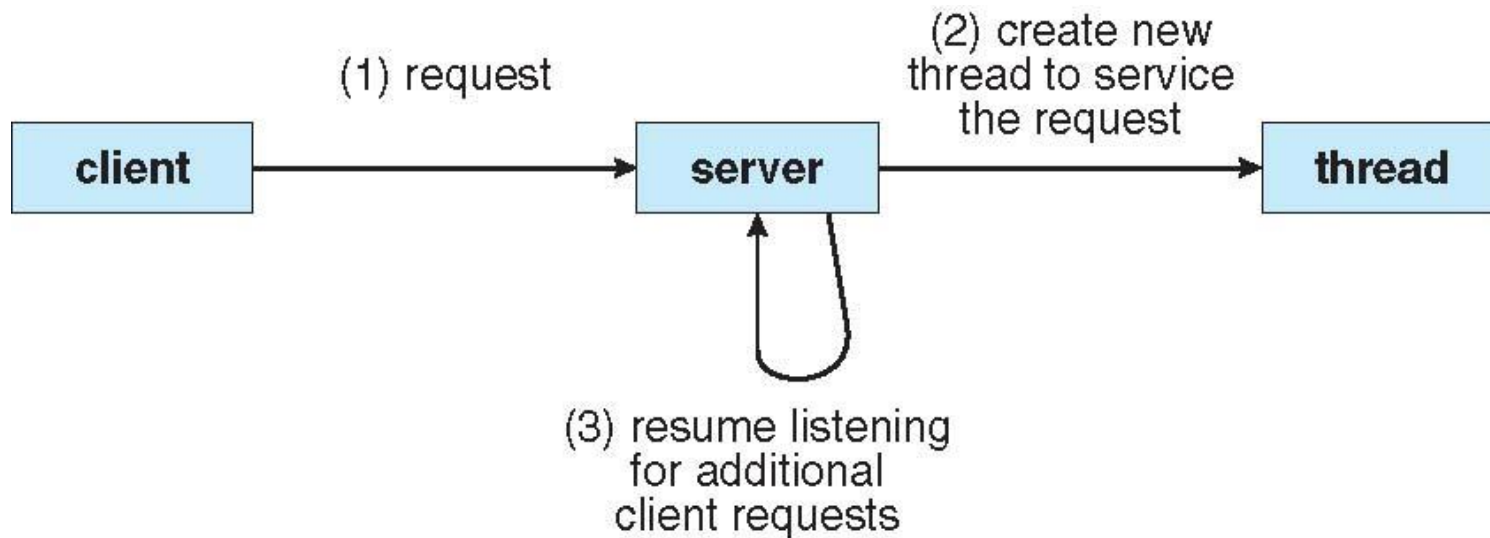


single-threaded process



multithreaded process

Multithreaded Server Architecture




Benefits

- Responsiveness
 - Single-threaded application will be unresponsive to user until a lengthy operations is completed
- Resource Sharing
 - Threads share memory and resources
- Economy
 - Thread management is less time consuming than process management
- Scalability
 - Multi-threaded process run on multi-core CPU


Multicore Programming

- Multicore systems putting pressure on programmers, challenges include
 - **Dividing activities**
 - **Balance**
 - **Data splitting**
 - **Data dependency**
 - **Testing and debugging**

Concurrent Execution on a Single-core System



Parallel Execution on a Multicore System




Multi-Threading Model

- User Thread
 - Thread management done by user-level threads library
 - Example: POSIX thread
- Kernel Thread
 - Supported by the Kernel
 - Examples: Windows, Linux, Mac OS X, etc.
- Multi-threading models
 - Many-to-One
 - One-to-One
 - Many-to-Many


Many-to-One Model

- Many user-level threads mapped to single kernel thread
- Discontinued because of inability to use multi-core
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads



One-to-one Model


- Each user-level thread maps to kernel thread
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later



Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows NT/2000 with the *ThreadFiber* package

Many-to-Many Model



Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS
- Example: POSIX thread - Pthread

Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard API for thread creation and synchronization
 - The Portable Operating System Interface (POSIX) is a family of standards specified by the IEEE Computer Society for maintaining compatibility between operating systems.
 - POSIX defines the application programming interface (API), along with command line shells and utility interfaces, for software compatibility with variants of Unix and other operating systems.
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)

Pthread Programming

```
/* thread.c */
#include <pthread.h>
#include <stdio.h>

int sum = 0; //shared by all threads
void *runner(void *param) {
    int i, upper = atoi(param);
    for (i=1;i<=upper;i++)
        sum += i;
    pthread_exit(0);
}
int main() {
    pthread_t tid; // thread identifier
    pthread_attr_t attr; //set of thread attributes
    pthread_attr_init(&attr); //get the default attributes
    /* create the thread */
    pthread_create(&tid, &attr, runner, "10");
    /* wait for the thread to exit */
    pthread_join(tid, NULL);
    printf("sum = %d\n", sum);
}

/* gcc -o test thread.c -lpthread */
```


Pthread Programming (Cont)

```
/* thread2.c */
#include <pthread.h>
#include <stdio.h>

int sum = 0; //shared by all threads

int main() {
    int i;
    pthread_t tid[10]; // thread identifier
    pthread_attr_t attr; //set of thread attributes
    pthread_attr_init(&attr); //get the default attributes
    /* create the thread */
    for(i=0; i < 10; i++)
        pthread_create(&tid[i], &attr, runner, "10");
    /* wait for the thread to exit */
    for(i=0; i < 10; i++)
        pthread_join(tid[i], NULL);
    printf("sum = %d\n", sum);
}
/* gcc -o test2 thread2.c -lpthread */
```

Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Thread cancellation of target thread
 - Asynchronous or deferred
- Signal handling
- Thread pools
- Thread-local storage

Semantics of `fork()` and `exec()`

- Does **`fork()`** duplicate only the calling thread or all threads?
- Some UNIX systems have chosen to have two versions of `fork()`
 - One version duplicates all threads
 - One version duplicates only the calling thread
- `exec()` semantics
 - Replace the entire process with the new program
- Which `fork()` to use is application dependent
 - If `exec()` is to be called: duplicate only the calling thread
 - If `exec()` is not to be called: duplicate all thread

Thread Cancellation

- Terminating a thread before it has finished
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- Challenges:
 - One thread is canceled while updating shared data
 - Asynchronous cancellation may cause problems
 - Resources have been allocated to a canceled thread
 - OS will reclaim system resource from a canceled thread but not all resources
 - System-wide resource may not be freed in asynchronous cancellation

Signal Handling

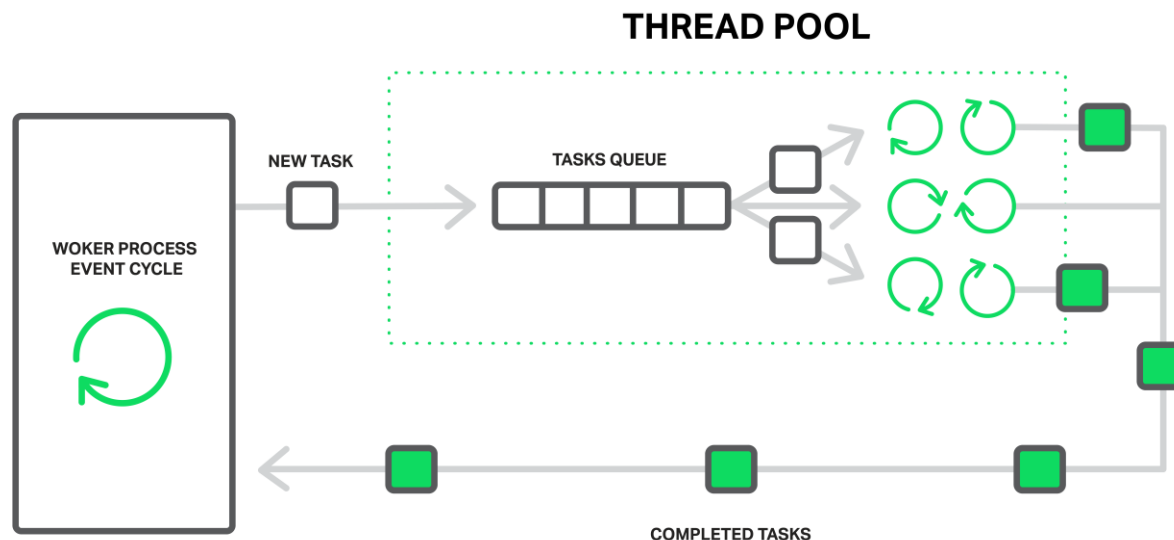
- Signals are used in UNIX systems to notify a process that a particular event has occurred
- A **signal handler** is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled
- A signal can be handled by
 - A default signal handler
 - A user-defined signal handler

Signal Handling (Cont)

- Synchronous signals and asynchronous signals
 - Synchronous signals: e.g., divided by 0
 - Asynchronous signals: e.g., <control><C>
- Handling signals in multi-threaded programs:
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process
- Signal handling:
 - Synchronous signals: deliver to the thread causing the signal
 - Asynchronous signals: tricky
 - <control><C> should be delivered to all

Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool



Thread Local Storage (TLS)

- Data is shared among threads by default
- TLS allows each thread to have its own copy of data
- Example
 - Thread identifier

End of Chapter 4